

Handwritten Digit Classification Using CNN and MNIST Dataset

Sora Owada

Mar 20, 2025

1. Describe

I am building a Convolutional Neural Network (CNN) that classifies handwritten digits (0–9) using the MNIST dataset. The project uses Python and TensorFlow to preprocess data, train a CNN model, and evaluate its performance. Currently, the model architecture and training process are implemented and functioning correctly in a Python script.

2. Examples

• **Input:** A grayscale image of a digit from the MNIST dataset (28x28 pixels).

Start Condition: The image is loaded, normalized, and reshaped into the input format of the CNN.

Output: The trained CNN model predicts a digit label (e.g., 3, 7, 0).

3. Meet the Objectives

• **Apply neural network architecture:**

I use TensorFlow's Keras API to define and train a CNN with multiple layers such as Conv2D, MaxPooling2D, and Dense layers.

• **Preprocess data:**

The program normalizes pixel values and reshapes the dataset into formats suitable for the CNN model.

• **Train and evaluate:**

I train the model using the training set and evaluate accuracy using the test set, tracking loss and performance.

• **Visualize performance:**

The code includes functions to display accuracy/loss curves and evaluate classification metrics.

4. Resources, Skills, Models, Tools, etc.

- **TensorFlow/Keras:**

Used to construct and train the CNN model. Offers high-level abstractions for layers and training.

- **MNIST dataset (via tensorflow.keras.datasets):**

Standard benchmark for digit recognition. Easily loadable for training/testing.

- **Matplotlib:** Used to visualize training/validation accuracy and loss over epochs.

- **Jupyter Notebook/Python Scripts:**

Tools for rapid development and debugging.

- **Online Tutorials (e.g., TensorFlow CNN tutorial):**

Helped in setting up and debugging model architecture.

5. Your Approach

Workflow Design:

1. **Input:** Load MNIST dataset using `tf.keras.datasets.mnist.load_data()`.

2. **Preprocessing:**

- Normalize pixel values to range `[0, 1]`.
- Reshape input images to `(28, 28, 1)`.
- One-hot encode the labels for training.

3. **Model Architecture:**

• `Conv2D -> MaxPooling2D -> Conv2D -> MaxPooling2D -> Flatten -> Dense -> Dense(10)`.

4. **Training:**

- Compile the model using Adam optimizer and sparse categorical cross-entropy.
- Train using `.fit()` for a set number of epochs with validation split.

5. **Evaluation:**

- Evaluate performance using `.evaluate()` on the test set.
- Plot training history (loss/accuracy).

6. **Output:** Predicted digit label with accuracy and loss reported.

6. Feedback

It would be helpful to receive feedback on:

- Whether the model complexity is appropriate for this task.
- Any best practices for structuring the CNN to improve accuracy.
- Suggestions on hyperparameter tuning or training visualization enhancements.
- Ideas for expanding the project (e.g., real-time digit prediction from user-drawn input).

7. References

MNIST Dataset: <https://www.kaggle.com/datasets/hojjatk/mnist-dataset>

GitHub: <https://github.com/gursky1/MNIST-Tensorflow-2>